


Identical architecture; measured separately with 1000 repeats

TCML CPU forward latency



TCML CUDA forward latency

